



# Artificial Neural Networks

Dr. Ahmad Altarawneh

Department of Data Science



Faculty of information technology, Mutah  
University

# Outline

- ▶ What are neural networks
- ▶ Perceptron: recalling
- ▶ Gradient descent
- ▶ Multilayer perceptron (MLP)
  - ▶ Mathematical representation
  - ▶ Forward pass
  - ▶ Calculating the error: loss function
- ▶ Backpropagation with gradient descent
- ▶ MLP implementation in Python (from scratch example)
  - ▶ Linear and nonlinear decision boundary
  - ▶ MLP for regression problems

# What are Neural Networks

- ▶ A Neural Network is a computational model inspired by the structure and functioning of the human brain.
- ▶ It consists of interconnected layers of nodes (called neurons), where each node processes input data, applies a mathematical function, and passes the output to the next layer.
- ▶ Nowadays, neural networks are used to identify patterns, make decisions, and solve complex problems in fields such as image recognition, natural language processing, and artificial intelligence.



# Perceptron

- ▶ A perceptron is the simplest NN, as it consists of a single neuron that processes (linear combinations) the input data, passes it through an activation function (step function), and produces an output
  - ▶ The perceptron classifier provides a linear decision boundary
- ▶ let a perceptron have  $n$  input features,  $x_1, x_2, \dots, x_n$  and corresponding weights  $w_1, w_2, \dots, w_n$

The perceptron calculates the Weighted Sum (Linear Combination):

$$z = x_1w_1 + x_2w_2 + \dots + x_nw_n + b$$

- ▶ The output of the linear combinations then enters to activation function, typically called a step function:




$$\hat{y} = f(z) = \begin{cases} 1, & \text{if } z \geq 0 \\ 0, & \text{if } z < 0 \end{cases}$$

Therefore

$$\hat{y} = f\left(\sum_{i=1}^n w_i x_i + b\right) = f(w \cdot x + b)$$


# Perceptron Cont.

- ▶  $w$  is called the weights vector, which defines the orientation of the decision boundary, and  $b$  is called bias, and controls its shift



- ▶ sometimes instead of a step function, a sigmoid function is used for better learning

$$\hat{y} = \frac{1}{1 + e^{-(z)}} = \frac{1}{1 + e^{-(w \cdot x + b)}}$$



# Perceptron Cont.

- ▶ After activating the weighted sum using the sigmoid function we evaluate the error
  - ▶ How much  $\hat{y}$  is far from a target value (actual class)  $y$
- ▶ This can be calculated using a function called the loss function
  - ▶ There are many loss functions
- ▶ The one we use here is called mean squared error (MSE), which is given by

$$error = (\hat{y} - y)^2$$

- ▶ The aim is to make this error as low as possible
  - ▶ low error rate means better learning
  - ▶ in classification tasks, 0 error means the line (perceptron) is classifying the data points without errors
- ▶ But how to reduce the error?

# Gradient decent

- ▶ We need to change the weights and bias (increase or decrease) in a way that makes the perceptron classify the data correctly
- ▶ But how do we know which weight to change, in which direction (increase or decrease), and the quantity of change?
  - ▶ We use a method called gradient descent
- ▶ gradient descent is a method based on differential equations (to find the slope of a function at a given point)
- ▶ There are different types of this based on how you use them
  - ▶ Batch gradient decent: calculate the gradients on the whole data, then update
  - ▶ Mini-batch gradient decent: divide the dataset into small batches, each batch contains number of examples (16, 32, 64 , etc.)
  - ▶ Stochastic gradient decent: update after every sample (our focus in the next slides)

# Gradient decent

- ▶ We need to change the weights and bias (increase or decrease) in a way that makes the perceptron classify the data correctly
- ▶ But how do we know which weight to change, in which direction (increase or decrease), and the quantity of change?
  - ▶ We use a method called gradient descent
- ▶ gradient descent is a method based on differential equations (to find the slope of a function at a given point)
- ▶ It involves calculating the **partial derivative** of the loss function w.r.t the components of the  $w$  vector and  $b$

$$\frac{\partial \text{loss}}{\partial w}$$

$$\frac{\partial \text{loss}}{\partial b}$$

- ▶ it gives how to change the  $w$  vector and  $b$  value so that the loss value is minimized (which we want to be 0)



# Gradient decent

## Cont.

- Assume that we have a single datapoint  $x = 1$ ,  $b = 0$  and the  $w$  initialized randomly to be  $w = 3$ , and the true label  $y$  for this point is 1

- we have 1 input, so we have 1  $w$  value

$$\hat{y} = w.x + b = 3 * 1 + 0 = 3$$

- using this  $w$  and  $b$  the predicted value is 4

$$loss = (\hat{y} - y)^2 = (3 - 1)^2 = 4$$

- to see how to change  $w$  we calculate the derivative of loss w.r.t  $w$

- However, this is a compound function, we do not have direct access to  $w$


- function inside another function

$$((w.x + b) - y)^2$$

- We need Chain rule!

the derivative of  $f(g(x))$  is  $f'(g(x)) \cdot g'(x)$

$$\frac{\partial \text{loss}}{\partial w} = \frac{\partial \text{loss}}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial w}$$



# Gradient decent

## Weights update

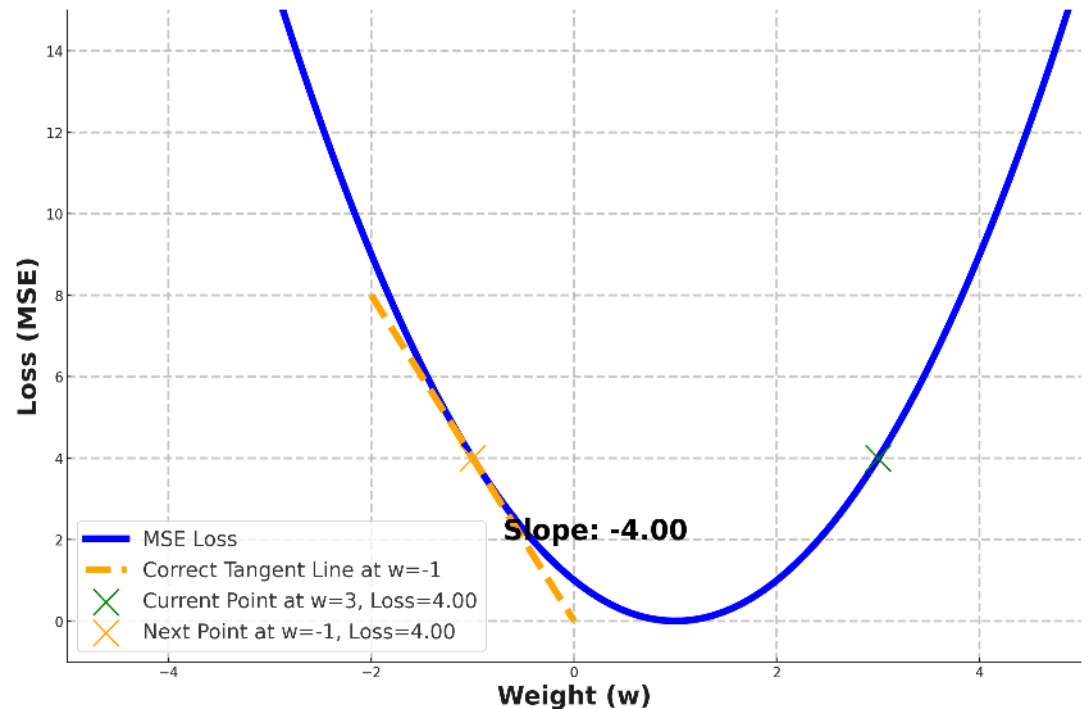
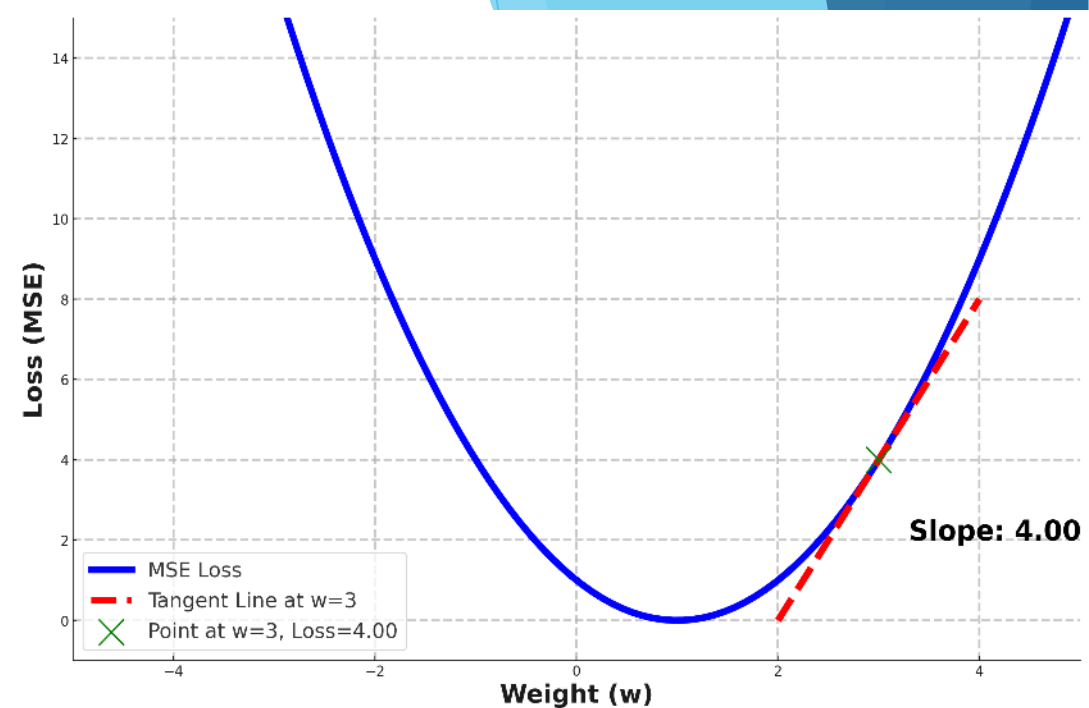
- ▶ Thus, the partial derivative of  $((w.x + b) - y)^2$  w.r.t  $w$  is:

$$\frac{\partial \text{loss}}{\partial w} = \frac{2 * (\hat{y} - y) * x}{\frac{\partial \text{loss}}{\partial \hat{y}}} \quad 2 * (3 - 1) * 1 = 4$$

- ▶ Therefore, the slope  $\frac{\partial \text{loss}}{\partial w}$  (technically called gradient) is 4
- ▶ now we change the  $w$ 
$$w_{\text{new}} = w_{\text{old}} - \frac{\partial \text{loss}}{\partial w} = 3 - 4 = -1$$
- ▶ Repeat until the function reaches its minimum.
- ▶ if we use sigmoid we add the derivative of it to the formula to become

$$2 * (\hat{y} - y) * \boxed{\text{sig}(z) * 1 - \text{sig}(z)} * x$$

where  $z = w.x + b$




# Perceptron

## Update weights with learning rate $\alpha$

- ▶ Note that, repeating in the current formula will keep the  $w$  fluctuating between 3 and  $-1$  forever, and never reaches the minimum
  - ▶ Therefore, we add a small value called  $\alpha$  to control the decent of  $w$
  - ▶  $\alpha$  is a value between 0-1, e.g., 0.01

$$w_{new} = w_{old} - \alpha \frac{\partial \text{loss}}{\partial w}$$

- ▶ The following figure shows several updates using  $\alpha = 0.8$
- ▶ if  $x$  has more than one input, and therefore  $w$ , we do the same w.r.t each component of  $w$



# Perceptron

## Bias update

- ▶ Note that we did not change  $b$  due to the simplicity of our example
- ▶ However,  $b$  needs to be changed in the same way, using gradient decent

$$\frac{\partial \text{loss}}{\partial b} = \frac{\partial \text{loss}}{\partial \hat{y}} * \frac{\partial \hat{y}}{\partial b}$$
$$\frac{\partial \text{loss}}{\partial b} = 2 * (\hat{y} - y) * 1$$

derivative of  $w \cdot x + b$  w.r.t  $b$

- ▶ and with sigmoid

$$\frac{\partial \text{loss}}{\partial b} = 2 * (\hat{y} - y) * \text{sig}(z) * 1 - \text{sig}(z) * 1$$

# Perceptron

## Python code

```
class Perceptron:
    def __init__(self, ninputs, epochs, alpha=0.01):
        self.n_epochs = epochs
        self.ninputs = ninputs
        self.weights = np.random.randn(ninputs)
        self.bias = 0
        self.alpha = alpha
        self.loss_history = []

    def sigmoid(self, x):
        return 1 / (1 + np.exp(-x))

    def mse(self, preds, y):
        return (preds - y) ** 2

    def dmse(self, preds, y):
        return 2 * (preds - y)

    def fit(self, X, y):
        for epoch in range(self.n_epochs):
            error = 0
            for sample, label in zip(X, y):
                z = np.dot(sample, self.weights) + self.bias
                preds = self.sigmoid(z)

                error += self.mse(preds, label)



                gradientsW = self.dmse(preds, label) * preds * (1 - preds) * sample
                gradientsB = self.dmse(preds, label) * preds * (1 - preds)

                self.weights -= self.alpha * gradientsW
                self.bias -= self.alpha * gradientsB

            avg_error = error / len(X)
            self.loss_history.append(avg_error)
            print(f'Epoch {epoch + 1}/{self.n_epochs}, Error: {avg_error}')

    def predict(self, x):
        y_hat = np.dot(x, self.weights) + self.bias
        return (self.sigmoid(y_hat) >= 0.5).astype(int)
```


# Perceptron Visualization



# Perceptron


## Final note

- ▶ Perceptron is weak in classifying nonlinearly separable data, as it is a linear model
  - ▶ Think of the XOR problem
- ▶ Impossible to solve using single perceptron
  - ▶ Can you find a line that can separate the following data points?
- ▶ We need something more complex
  - ▶ Multilayer Perceptron (MLP)



# Multilayer Perceptron

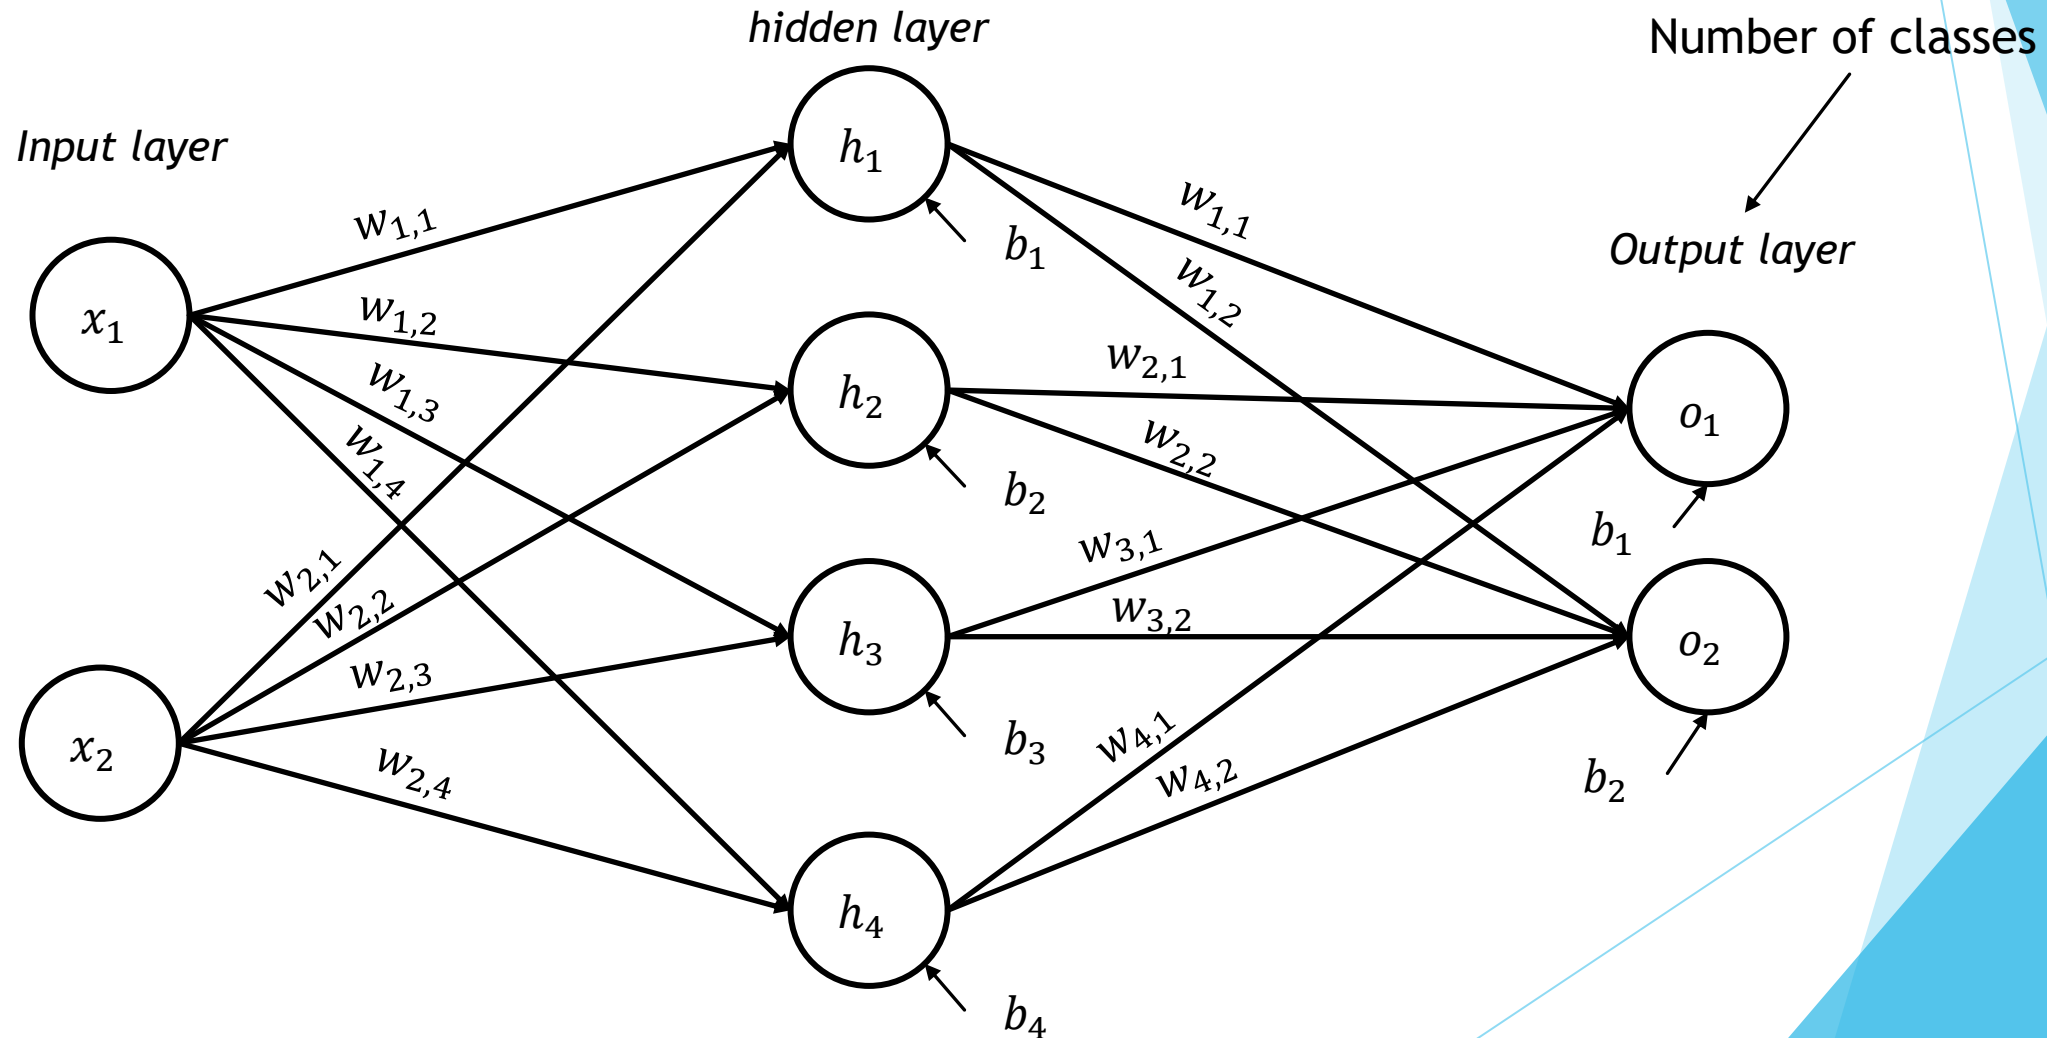
- ▶ An MLP is a type of artificial neural network that consists of multiple layers of Perceptron (neurons), including an input layer, one or more hidden layers, and an output layer
  - ▶ Input Layer: The first layer that receives the input features.
  - ▶ Hidden Layers: One or more intermediate layers where the actual processing is done through weighted connections. Each neuron in a hidden layer applies a non-linear activation function to the weighted sum of inputs.
  - ▶ Output Layer: The final layer that produces the output of the network, representing the result of the prediction or classification.





# Cont.

- ▶ The following present an example architecture of MLP
  - ▶ Each node in the hidden and output layers has associated weights and bias



# MLP, Cont.

- ▶ The training process of MLP network is similar to that of perceptron, using gradient decent
  - ▶ it consists of forward pass and then correct the weights based on the error in a step called backpropagation
- ▶ In the **forward pass** we calculate the linear combination of each single node in the network sequential order using the same formula used in the perceptron

$$\hat{y} = w \cdot x + b$$

- ▶ Then we calculate the error for each output node
- ▶ In the **backward pass** we correct the parameters  $w$  and  $b$  in whole architecture in a reversed sequential order

# Cont.

The forward pass of the previous architecture looks as follows

$$h_1 = x_1 * w_{1,1} + x_2 * w_{2,1} + b_1$$

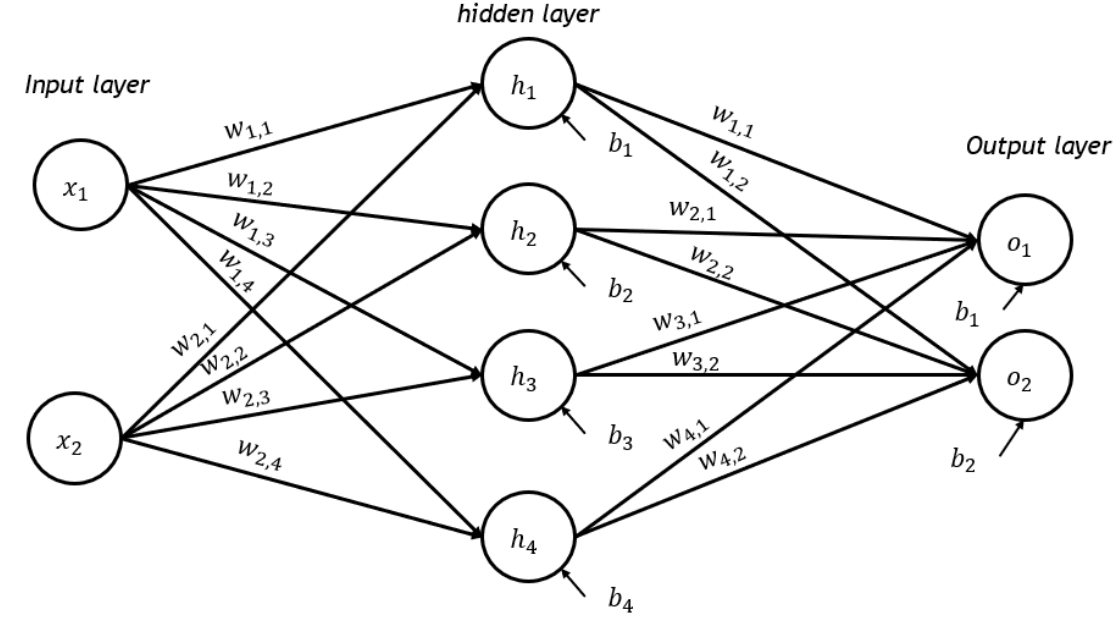
$$h_2 = x_1 * w_{1,2} + x_2 * w_{2,2} + b_2$$

$$h_3 = x_1 * w_{1,3} + x_2 * w_{2,3} + b_3$$

$$h_4 = x_1 * w_{1,4} + x_2 * w_{2,4} + b_4$$

$$o_1 = h_1 * w_{1,1} + h_2 * w_{2,1} + h_3 * w_{3,1} + h_4 * w_{4,1} + b_1$$


$$o_2 = h_1 * w_{1,2} + h_2 * w_{2,2} + h_3 * w_{3,2} + h_4 * w_{4,2} + b_2$$



But MLP network typically include much greater number of nodes, which makes the calculations, in the above manner, complex and inefficient

SOLUTION: Use matrix operations

# Cont. MLP Matrices



- ▶  $w$  in each layer can be represented as weight matrix of size  $m \times n$ , where  $m$  is the input values, and  $n$  is the output values (number of nodes in the next layer)

- ▶  $w$  in the first layer is  $2 \times 4$  ; 2 is the number of inputs to this layer ( $x_1$  and  $x_2$ ) and 4 is the number of outputs ( $h_1, h_2, h_3$  and  $h_4$ )

$$w1 = \begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} & w_{1,4} \\ w_{2,1} & w_{2,2} & w_{2,3} & w_{2,4} \end{bmatrix}$$

- ▶ Also, the inputs  $x$  can be represented as a matrix  $s \times m$ , where  $s$  is the number of samples
  - ▶ if we pass one sample at a time (Stochastic gradient decent), then  $s = 1$
  - ▶  $m$  is the number of values in  $x$

$$x = [x_1 \quad x_2]$$

- ▶ Also  $b$  in the hidden layer  $h$  is a vector matrix of size  $1 \times 4$

$$b = [b_1, b_2, b_3 \text{ and } b_4]$$

# Cont.

- ▶ Now instead of calculating the  $h$  values individually we perform matrix multiplication to directly calculate the terms in the hidden layer  $h_1, h_2, h_3$  and  $h_4$

$$h = x.w1 + b1$$

- ▶ remember the inner dimension should match in the matrix multiplications

$$h = [x_1 \quad x_2] \cdot \begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} & w_{1,4} \\ w_{2,1} & w_{2,2} & w_{2,3} & w_{2,4} \end{bmatrix} + [b_1 \quad b_2 \quad b_3 \quad b_4]$$

- ▶ matrix multiplication is performed by weight sum row from the first matrix and the column from the second matrix
- ▶ The output has the outer dimension of the multiplied matrices  $s \times n$ , in our case  $1 \times 4$

# Cont.

- ▶ This how it works

$$h_{temp} = \begin{bmatrix} x_1 & x_2 \end{bmatrix} \cdot \begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} & w_{1,4} \\ w_{2,1} & w_{2,2} & w_{2,3} & w_{2,4} \end{bmatrix}$$

$$h_{temp} = [x_1 * w_{1,1} + x_2 * w_{2,1}, x_1 * w_{1,2} + x_2 * w_{2,2}, x_1 * w_{1,3} + x_2 * w_{2,3}, x_1 * w_{1,4} + x_2 * w_{2,4}]$$

- ▶ Then we add, elementwise summation, the bias terms to get the final outputs of in  $h$

$$h = h_{temp} + b1$$

$$h = [x_1 * w_{1,1} + x_2 * w_{2,1} + b, x_1 * w_{1,2} + x_2 * w_{2,2} + b, x_1 * w_{1,3} + x_2 * w_{2,3} + b, x_1 * w_{1,4} + x_2 * w_{2,4} + b]$$

$$h = [h_1, h_2, h_3, h_4]$$

- ▶ Now this  $h$  becomes the input for the next layer and we repeat until we get the outputs
- ▶ Same results we get before using math operations, but what makes it efficient is that this process is equivalent to dot product and vector summation in numpy, Python

$$\mathbf{h} = \mathbf{np.dot(x, w) + b}$$

# Example

- ▶ Now let us take example to see how to calculate the forward pass using our architecture
  - ▶ Assume that our dataset has 1 example
  - ▶  $y$  (class) should be one-hot encoded, having 1 at index 0 means that the actual class is 0

$$x = [3 \quad 5] \quad y = [1 \quad 0]$$

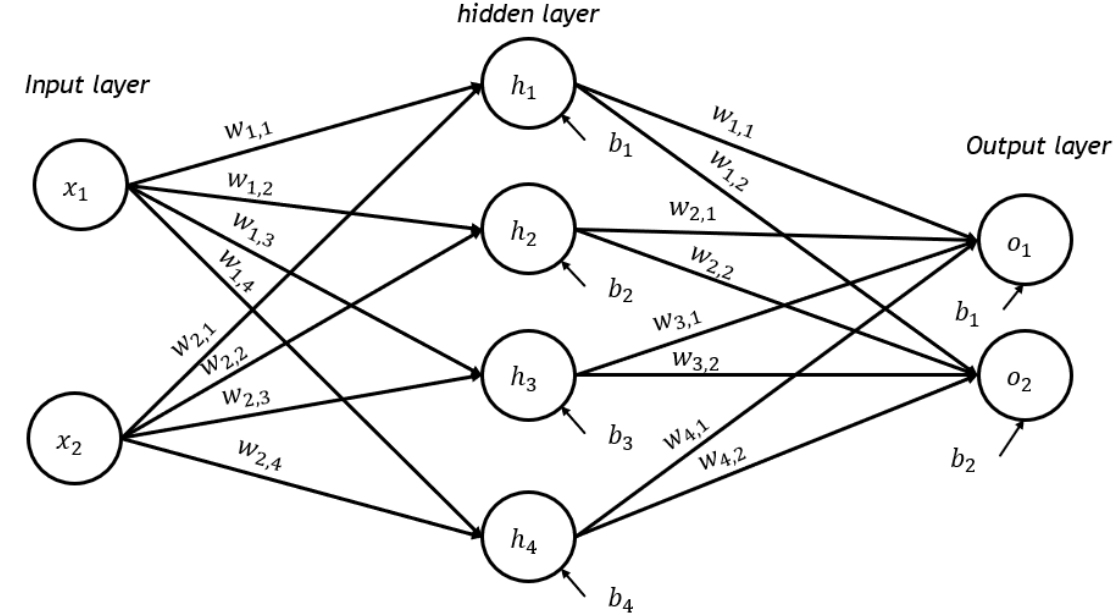
- ▶  $w$  is randomly initialized, and  $b$  is a vector of ones

$$x = [3 \quad 5] \quad w1 = \begin{bmatrix} 0.5 & 1 & 7 & 3 \\ 1 & 2 & 1 & 0.5 \end{bmatrix} \quad b1 = [1 \quad 1 \quad 1 \quad 1]$$

$$h = [3 * 0.5 + 5 * 1 + 1, \quad 3 * 1 + 5 * 2 + 1, \quad 3 * 7 + 5 * 1 + 1, \quad 3 * 3 + 5 * 0.5 + 1]$$
$$h = [7.5, 14, 27, 12.5]$$

```
x = np.array([3, 5])
w = np.array([[0.5, 1, 7, 3], [1, 2, 1, 0.5]])
b = np.array([1, 1, 1, 1])
h = np.dot(x, w) + b

array([ 7.5, 14. , 27. , 12.5])
```



# Example, Cont.

- ▶ Now  $h$  becomes the input to the new layer which has  $w_{4 \times 2}$

$$h = [7.5, 14, 27, 12.5] \quad w2 = \begin{bmatrix} 1 & 2 \\ 0.5 & 0.5 \\ 3 & 1 \\ 2 & 1 \end{bmatrix} \quad b2 = [1 \quad 1 \quad 1 \quad 1]$$


$$o = [7.5 * 1 + 14 * 0.5 + 27 * 3 + 12.5 * 2 + 1, \quad 7.5 * 2 + 14 * 0.5 + 27 * 1 + 12.5 * 1 + 1]$$

$$o = [121.5, 62.5]$$

- ▶ The error is calculated between the predicted and the actual using MSE loss
  - ▶ its equation now is a bit different as we have multi label

$$error = \frac{1}{n} \sum_{i=1}^n (o_i - y_i)^2$$

$$error = \frac{1}{2} \sum_{i=1}^n (o_i - y_i)^2 = \frac{1}{2} * (121.5 - 1)^2 + (62.5 - 0)^2 = \frac{1}{2} * 14,520 + 3,906 = 9,213$$





# Backward pass

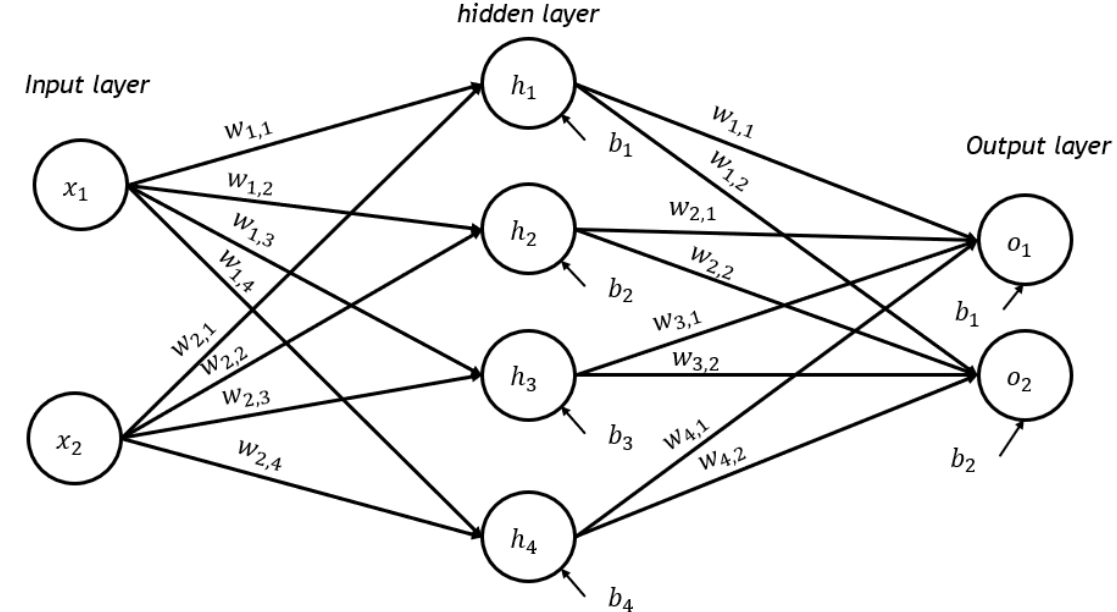
- ▶ Now we have the error, and we want to adjust the weights and the bias terms in both layers to reduce this error
  - ▶ Using backpropagation algorithm (gradient decent)
- ▶ remember from chain rule we have to calculate the derivatives w.r.t  $w_2, b_2, w_1$  and  $b_1$
- ▶ we start with calculating the gradient of loss w.r.t each element in  $o$ ,  $\frac{\partial \text{loss}}{\partial o_i}$ 
  - ▶ note here  $o$  is  $\hat{y}$  in the previous example (predictions)

$$\frac{\partial \text{loss}}{\partial o_i} = \frac{2}{2} (o_i - y_i)$$

$$\frac{\partial \text{loss}}{\partial o_1} = 1 * (121.5 - 1) = 120.5$$

$$\frac{\partial \text{loss}}{\partial o_2} = 1 * (62.5 - 0) = 62.5$$

$$\frac{\partial \text{loss}}{\partial o} = [120.5, 62.5]$$



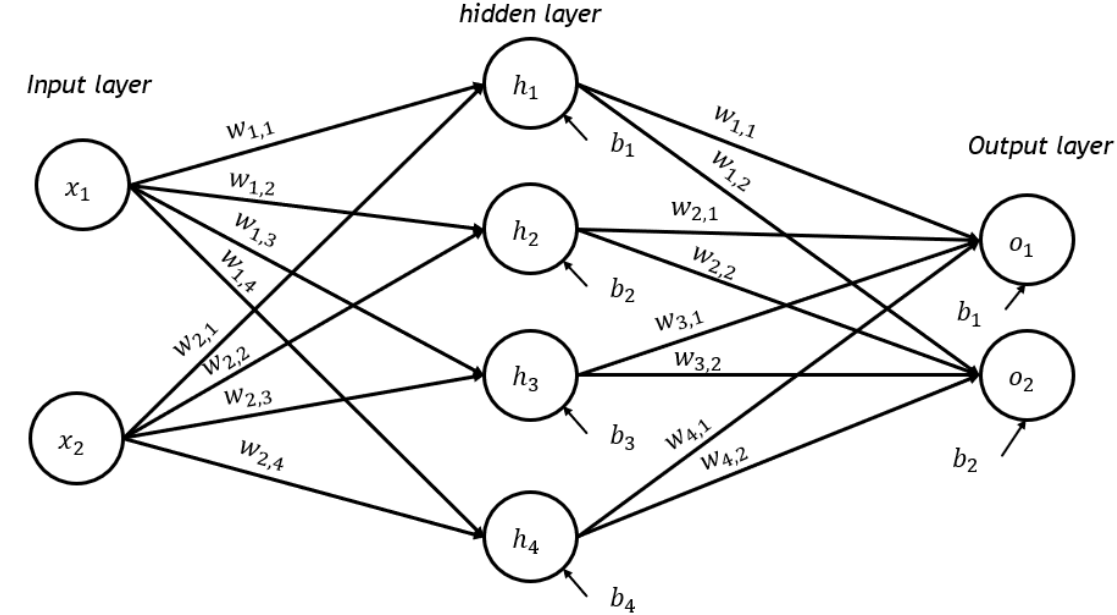
# Backward pass, Cont.

- Now we have to find the derivative w.r.t  $w_2$  and  $b_2$

$$\frac{\partial \text{loss}}{\partial w_2} = \begin{bmatrix} \frac{\partial \text{loss}}{\partial w_{1,1}} & \frac{\partial \text{loss}}{\partial w_{1,2}} \\ \frac{\partial \text{loss}}{\partial w_{2,1}} & \frac{\partial \text{loss}}{\partial w_{2,2}} \\ \frac{\partial \text{loss}}{\partial w_{3,1}} & \frac{\partial \text{loss}}{\partial w_{3,2}} \\ \frac{\partial \text{loss}}{\partial w_{4,1}} & \frac{\partial \text{loss}}{\partial w_{4,2}} \end{bmatrix}$$

$$\frac{\partial \text{loss}}{\partial w_{1,1}} = \frac{\partial \text{loss}}{\partial o_1} * \frac{\partial o_1}{\partial w_{1,1}} = 120.5 * 7.5 = 903.75$$

$$\frac{\partial \text{loss}}{\partial w_{1,2}} = \frac{\partial \text{loss}}{\partial o_2} * \frac{\partial o_2}{\partial w_{1,2}} = 62.5 * 7.5 = 468.75$$



$$\frac{\partial o_1}{\partial w_{1,1}} = h_1$$

$$\frac{\partial o_1}{\partial w_{1,2}} = h_2$$

# Cont.

- ▶ This calculations can be done simply using matrix operations

$$\frac{\partial \text{loss}}{\partial w2} = \begin{bmatrix} 7.5 \\ 14 \\ 27 \\ 12.5 \end{bmatrix} \cdot [121.5, 62.5]$$


$$h^T \cdot o \quad \frac{\partial \text{loss}}{\partial w2} = \begin{bmatrix} 903.75 & 468.75 \\ 1687 & 875 \\ 3253.5 & 1687.5 \\ 1506.25 & 781.25 \end{bmatrix}$$

$$\frac{\partial \text{loss}}{\partial b2} = \frac{\partial \text{loss}}{\partial o} * \frac{\partial o}{\partial b2} = [121.5, 62.5] * [1, 1] = [121.5, 62.5]$$

- ▶ Although these numbers looks large, we update the weights in  $w2$  with a proportion of them
  - ▶ do not forget the learning rate  $\alpha$


$$w2_{\text{new}} = w2_{\text{old}} - \alpha * \frac{\partial \text{loss}}{\partial w2}$$

$$b2_{\text{new}} = b2_{\text{old}} - \alpha * \frac{\partial \text{loss}}{\partial b2}$$



# Cont.

- ▶ In this way we successfully updated the weights  $w_2$  and bias  $b_2$  in the second layer
  - ▶ But what about the weights and bias in the first layer,  $w_1$  and  $b_1$ ?



- ▶ for this we need to calculate the gradient with respect to the input  $\frac{\partial \text{loss}}{\partial h}$ , so that we pass it to the previous layer and repeated the calculations
  - ▶ it is simply given by

$$\frac{\partial \text{loss}}{\partial h} = [121.5, 62.5] \cdot w_{2_{old}}^T$$

- ▶ The above dot product gives a matrix  $1 \times 4$  which is the gradient w.r.t  $h$ 
  - ▶ Use it to calculate for the  $w_1$  and  $b_1$

# Cont.

- ▶ Note that in the previous explanation we did not include any activation functions for **simplicity**
- ▶ The MLP will remain linear no matter how layers or nodes you add
- ▶ to solve this, we use a nonlinear activation function to activate the output of each node in the hidden or output layers
  - ▶ Sigmoid for example
- ▶ the math will work in the same way with additional step, which is calculating and propagating the derivative of sigmoid to the process

# Code

- ▶ Let us see how can we implement this architecture and all its math using Python
- ▶ We implement a class called layer, so that each object (layer) has its own weights and bias stored in it and perform the operations independently

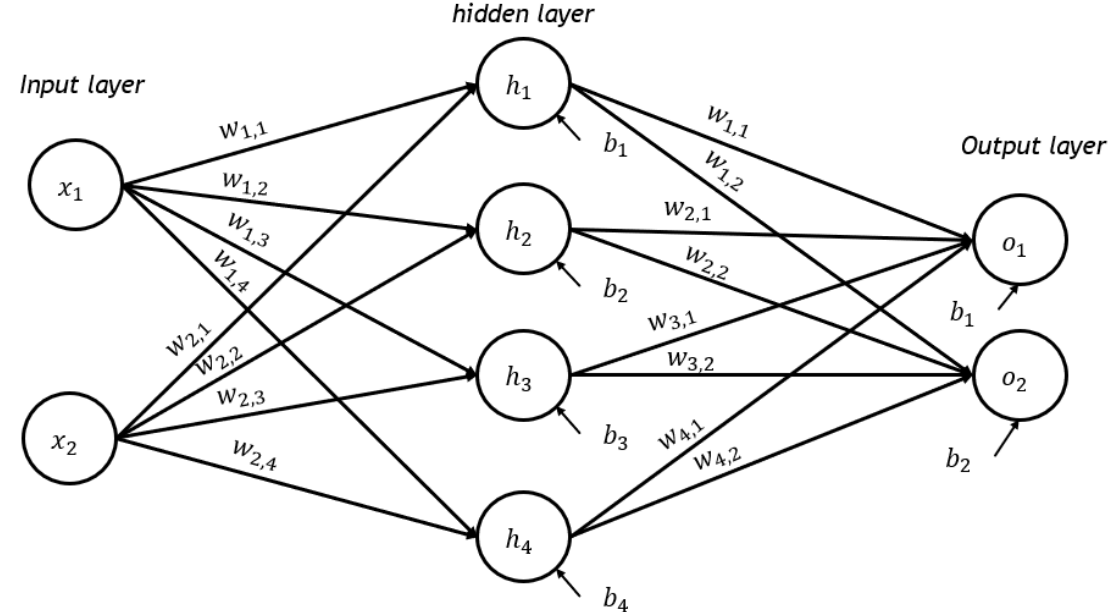
```
import numpy as np

class Dense:
    def __init__(self, inF, outF):
        self.weights = np.random.randn(inF, outF)
        self.biases = np.ones((1, outF))

    def forward(self, inputs):
        self.inputs = inputs
        return np.dot(inputs, self.weights) + self.biases

    def backward(self, dvalues):
        self.dweights = np.dot(self.inputs.T, dvalues)
        self.dbiases = np.sum(dvalues, axis=0, keepdims=True)
        self.dinputs = np.dot(dvalues, self.weights.T)

layer1 = Dense(2, 4)
layer2 = Dense(4, 2)
```



Let us check if we use the data in our example, we will get the same results

# Code, Cont.

```
layer1.weights = np.array([[0.5, 1, 7, 3],[1, 2, 1, 0.5]])
layer2.weights = np.array([[1, 2],[0.5, 0.5],[3, 1],[2, 1]])

x = np.array([[3,5]])
y = np.array([[1,0]])

h = layer1.forward(x)
print(f"the output of the first layer is {h}")
o = layer2.forward(np.array(h))
print(f"the output of the second (output) layer is {o}")

def mse_loss(y_true, y_pred):
    return np.mean((y_pred - y_true) ** 2)

# Derivative of MSE Loss
def mse_loss_derivative(y_true, y_pred):
    return 2 * (y_pred - y_true) / y_true.size

loss = mse_loss(y, o)
print(f"Mean Squared Error Loss: {loss}")

dL_do = mse_loss_derivative(y, o)
layer2.backward(dL_do)
layer1.backward(layer2.dinputs)

print(f"Gradients of layer2 weights: \n {layer2.dweights}")
print(f"Gradients of layer2 biases: \n {layer2.dbiases}")

print(f"Gradients of layer1 weights: \n {layer1.dweights}")
print(f"Gradients of layer2 biases: \n {layer2.dbiases}")
```

the output of the first layer is  $\begin{bmatrix} 7.5 & 14. & 27. & 12.5 \end{bmatrix}$   
the output of the second (output) layer is  $\begin{bmatrix} 121.5 & 62.5 \end{bmatrix}$   
Mean Squared Error Loss: 9213.25  
Gradients of layer2 weights:  
 $\begin{bmatrix} 903.75 & 468.75 \\ 1687. & 875. \\ 3253.5 & 1687.5 \\ 1506.25 & 781.25 \end{bmatrix}$   
Gradients of layer2 biases:  
 $\begin{bmatrix} 120.5 & 62.5 \end{bmatrix}$   
Gradients of layer1 weights:  
 $\begin{bmatrix} 736.5 & 274.5 & 1272. & 910.5 \\ 1227.5 & 457.5 & 2120. & 1517.5 \end{bmatrix}$   
Gradients of layer2 biases:  
 $\begin{bmatrix} 120.5 & 62.5 \end{bmatrix}$

# Code, Cont.

- Update the weights and biases in these layers and recheck, using the same point, if the loss is decreasing

```
alpha = 0.001

layer1.weights -= alpha * layer1.dweights
layer2.weights -= alpha * layer2.dweights

layer1.biases -= alpha * layer1.dbiases
layer2.biases -= alpha * layer2.dbiases

h = layer1.forward(x)
o = layer2.forward(h)
loss = mse_loss(y, o)
print(f"Mean Squared Error Loss: {loss}")
```

Loss: 196.56

**The loss decreases.  
There is learning!**



## Example using dataset

```
X, y = make_blobs(n_samples=100, centers=2, n_features=2, random_state=42)



one_hot_encoder = OneHotEncoder(sparse=False)
y_onehot = one_hot_encoder.fit_transform(y.reshape(-1, 1))  # Shape: (100, 2)

layer1 = Dense(2, 4)
layer2 = Dense(4, 2)

alpha = 0.01
epochs = 10
# Training loop
for epoch in range(epochs):
    loss_epoch = 0
    for i in range(len(X)):
        x_sample = np.expand_dims(X[i], axis=0)  # (shape: (1, 2))
        y_sample = np.expand_dims(y_onehot[i], axis=0)  # (shape: (1, 2))

        h = layer1.forward(x_sample)
        o = layer2.forward(h)  # Raw output
        # loss
        loss = mse_loss(y_sample, o)
        loss_epoch += loss
        # Calculate gradients
        dL_do = mse_loss_derivative(y_sample, o)
        layer2.backward(dL_do)
        layer1.backward(layer2.dinputs)
        # Update w and b
        layer1.weights -= alpha * layer1.dweights
        layer1.biases -= alpha * layer1.dbiases
        layer2.weights -= alpha * layer2.dweights
        layer2.biases -= alpha * layer2.dbiases

    # Print average loss for the epoch
    if epoch % 1 == 0:
        print(f"Epoch {epoch}, Loss: {loss_epoch / len(X)}")
```



# Activation functions


- ▶ For the previous example we note that the form of MLP we used before cannot separate nonlinearly separable
- ▶ to make it able to solve this kind of data we have to add non linear activation functions to the MLP after each layer, as we discussed earlier
- ▶ There are many activation functions
  1. Sigmoid (discussed before)
  2. Tanh
  3. ReLU
  4. Leaky ReLU
  5. Softmax

# Tanh: the Hyperbolic function

- ▶ Tanh is similar to the Sigmoid except that it squishes the values between -1 and 1 instead of 0 and 1
- ▶ it is given by the following expression

$$\tanh = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$


$$\tanh' = 1 - \tanh^2(x)$$



# ReLU: Rectified Linear Unit

- ▶ ReLU is the most common activation function in deep learning as it helps preventing what is called **gradient vanishing** problem
  - ▶ Gradients become so small and eventually become zeros

$$\text{ReLU}(x) = \begin{cases} 0, & \text{if } x < 0 \\ x, & \text{if } x \geq 0 \end{cases} \quad \text{ReLU}'(x) = \begin{cases} 0, & \text{if } x < 0 \\ 1, & \text{if } x \geq 0 \end{cases}$$




# leaky ReLU

- ▶ the standard ReLU suffers from a problem called dead neuron dying ReLU
  - ▶ this happens when the input to ReLU is negative
- ▶ So instead of zeroing negative elements we use a small scaler for negative values

$$\text{ReLU}(x) = \begin{cases} \alpha x, & \text{if } x < 0 \\ x, & \text{if } x \geq 0 \end{cases}$$

$$\text{ReLU}'(x) = \begin{cases} \alpha, & \text{if } x < 0 \\ 1, & \text{if } x \geq 0 \end{cases}$$



# Softmax

- ▶ Softmax is another activation function that is typically used on the output layer of multiclass classification tasks.
- ▶ It transforms the raw output logits from the network into a **probability distribution**, where each class is assigned a probability, and the sum of probabilities across all classes equals 1
  - ▶ For a problem of  $n$  classes, the Softmax is given by

$$\text{Softmax}(x_i) = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$

- ▶ The derivative of Softmax is tricky, will discuss it later on. For this reason usually the softmax is combined with a type of loss called **cross-entropy loss** to simplify its derivative

# Python code

► Now let us add this activation function to our implementation and see how the decision boundary looks

```
class ActivationRelu:
    def forward(self, inputs):
        self.inputs = inputs
        self.output = np.maximum(0, inputs)
        return self.output

    def backward(self, dvalues):
        self.dinputs = dvalues * np.where(self.inputs > 0, 1, 0)

class ActivationSig:
    def forward(self, inputs):
        self.inputs = inputs
        self.output = 1 / (1 + np.exp(-inputs))
        return self.output

    def backward(self, dvalues):
        # Derivative of Sigmoid function
        self.dinputs = dvalues * (self.output * (1 - self.output))
```

► Also, we need to increase the number of layers and nodes to make the MLP able to fit more complex data

# Example with ReLU

```
layer1 = Dense(2, 10)
activation1 = ActivationRelu()
```


```
layer2 = Dense(10, 10)
activation2 = ActivationRelu()
```

```
layer3 = Dense(10, 2)
activation3 = ActivationRelu()
```

```
# Training parameters
```

```
alpha = 0.1
```

```
epochs = 1000
```



```
# Training loop
for epoch in range(epochs):
    loss_epoch = 0
    for i in range(len(X)):
        # Get one sample and one-hot label
        x_sample = np.expand_dims(X[i], axis=0)
        y_sample = np.expand_dims(y_onehot[i], axis=0)

        # Forward pass
        h1 = layer1.forward(x_sample)
        h1 = activation1.forward(h1)
        h2 = layer2.forward(h1)
        h2 = activation2.forward(h2)
        o = layer3.forward(h2)
        o = activation3.forward(o)
        # Compute loss
        loss = mse_loss(y_sample, o)
        loss_epoch += loss

        # Backward pass
        dL_do = mse_loss_derivative(y_sample, o)
        activation3.backward(dL_do)
        layer3.backward(activation3.dinputs)
        activation2.backward(layer3.dinputs)
        layer2.backward(activation2.dinputs)
        activation1.backward(layer2.dinputs)
        layer1.backward(activation1.dinputs)
        # Update weights and biases
        layer1.weights -= alpha * layer1.dweights
        layer1.biases -= alpha * layer1.dbiases
        layer2.weights -= alpha * layer2.dweights
        layer2.biases -= alpha * layer2.dbiases
        layer3.weights -= alpha * layer3.dweights
        layer3.biases -= alpha * layer3.dbiases


    # Print average loss for the epoch
    if epoch % 10 == 0:
        print(f"Epoch {epoch}, Loss: {loss_epoch / len(X)}")
```




# Visualization code

```
def plot_decision_boundary():  
    # Define a grid for the contour plot  
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1  
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1  
    xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.1),  
                          np.arange(y_min, y_max, 0.1))  
  
    grid_points = np.c_[xx.ravel(), yy.ravel()]  
    h1 = layer1.forward(grid_points)  
    h1 = activation1.forward(h1)  
  
    h2 = layer2.forward(h1)  
    h2 = activation2.forward(h2)  
  
    o = layer3.forward(h2)  
  
    Z = np.argmax(o, axis=1).reshape(xx.shape)  
  
    plt.contourf(xx, yy, Z, levels=25, cmap="RdYlBu", alpha=0.8)  
    plt.scatter(X[:, 0], X[:, 1], c=y, s=40, cmap="RdYlBu", edgecolor="k")  
    plt.title("Decision Boundary with 3-Layer Network")  
    plt.xlabel("Feature 1")  
    plt.ylabel("Feature 2")  
    plt.show()
```

## Using ReLU



## Using Softmax



# Loss functions

- ▶ The loss function we were using so far is the MSE loss, which is usually works for regression tasks
  - ▶ Although it can work for classification tasks
- ▶ There are better loss functions designed for classification tasks, such as
  - ▶ MSE: Regression
  - ▶ MAE: Regression
  - ▶ binary cross entropy: two classes
  - ▶ categorical cross entropy: labels are one-hot-encoded integers
  - ▶ sparse categorical cross entropy: labels are integers
  - ▶ and many more ...

# Cross entropy

- ▶ for a classification problem with  $y$  one-hot-encoded, The formula for **Cross Entropy Loss** is

$$CELoss = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^C y_{i,j} \log(\hat{y}_{i,j})$$

where:

- ▶  $N$  is the number of examples
- ▶  $C$  is the number of classes
- ▶  $y_{i,j}$  is the true label for sample  $i$  and class  $j$ , where  $y_{i,j} = 1$  if the class is the true class for the sample, and 0 otherwise.
- ▶  $\hat{y}_{i,j}$  is the predicted probability (or softmax output) for sample  $i$  and class  $j$ .
- ▶ For binary classification

$$\text{Loss} = -\frac{1}{N} \sum_{i=1}^N [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

# Categorical cross entropy with Softmax

- ▶ Derivative of cross entropy is tricky, but combining it with Softmax simplifies implementation

$$\frac{\partial \text{Loss}}{\partial z_i} = \hat{y}_i - y_i \quad i \text{ is the class in the one-hot vector}$$

```
class SoftmaxCrossEntropy:

    def forward(self, logits, y_true):
        exp_values = np.exp(logits - np.max(logits, axis=1, keepdims=True))
        self.y_pred = exp_values / np.sum(exp_values, axis=1, keepdims=True)

        self.y_pred = np.clip(self.y_pred, 1e-7, 1 - 1e-7)
        # Compute cross-entropy loss
        loss = -np.sum(y_true * np.log(self.y_pred)) / y_true.shape[0] # Average over batch
        return loss

    def backward(self, y_true):
        # Gradient of the loss w.r.t logits
        grad = self.y_pred - y_true
        return grad
```

`-np.max(logits))` is used to avoid very large and very low numbers